



INDIA.RUNS

Build what next India runs on

Team Name : (your team)

Team Leader Name :

Problem Statement : Rank the top 100 of 100,000 candidates for one nuanced Senior AI Engineer JD with recruiter-grade judgment — not keyword matching — avoiding keyword-stuffer traps and ~80 planted honeypots. Project: Lighthouse — a reasoning-first, explainable candidate ranker.

The Problem

- 100,000 candidates; genuine Senior-AI-Engineer fits are rare — only 0.78% even have an AI-ish title.
- Seeded traps: keyword-stuffers (an Accountant listing 9 AI skills), plain-language Tier-5s (real systems, no buzzwords), behavioral twins, and ~80 honeypots (subtly impossible profiles).
- The provided sample_submission ranks the stuffers #1–20 — the exact wrong answer. >10% honeypots in the top-100 = disqualification.
- Scoring is NDCG@10-heavy: the top-10 must be right. The real task is reasoning about the gap between what the JD says and what it means.

Solution Overview

- Lighthouse — an offline-LLM-augmented hybrid ranker: recruiter-grade, reasoning-first, fully explainable.
- Five fit components (semantic + role-coherence + career-evidence + experience + skill-trust), gated by JD hard-negatives and a behavioral modifier, with honeypots zeroed.
- Reads career evidence & semantics, not keywords — surfaces plain-language Tier-5s and sinks keyword-stuffers.
- Every pick ships a grounded, no-hallucination reason.
- CPU-only, < 5 min, no API at rank-time — production-shaped, not a GPT-per-candidate demo.
- Live Discovery surface: paste JD prose → auto-facets (or preset chips), then look-alike search across the 100K — recruiter-grade exploration, not a static list.

JD Understanding & Candidate Evaluation

- Must-haves: production embeddings/retrieval, vector/hybrid search ops, ranking-eval (NDCG/MRR/MAP), strong Python. Ideal 6–8 yrs, 4–5 applied-ML at product (not services) cos, shipped a ranking/search/recsys system.
- 8 hard-negative gates: services-only, outside-India & won't-relocate (no visa), research-only, CV/speech-only w/o NLP, LangChain-only-recent, title-chasers, non-engineering current role, unbacked-expertise (≥ 5 expert-skill claims absent from every role — the stuffer trap). Dampen, don't annihilate.
- Signal priority: career-history evidence > job title > listed skills. Behavioral signals decide actual availability.
- Claude (offline) parsed the JD into a static committed rubric — read at rank-time with no API call.

Ranking Methodology

- Retrieve: precomputed bge-small embeddings over clean per-candidate text blobs (100K). (A BM25 lexical index is built offline — optional, not loaded at rank-time.)
- Score, each in [0,1]: semantic_fit (cosine vs JD facets, top-4 blend), role_coherence (taxonomy+semantic), career_evidence (built ranking/search at product cos; recency-weighted, 4-yr half-life), experience_fit (soft 6–8 curve), trust_skills (proficiency×duration×endorsements×assessment, +AI/ML certs).
- Combine: weighted base (role_coherence 0.26 & career_evidence 0.24 lead) × $\prod(8 \text{ gates})$ × behavioral modifier [0.80–1.12] (activity, response-rate, recruiter-saves/views, open-to-work lift); honeypots → 0.
- final = base × gates × behavior. Order by score desc; ties → candidate_id ascending; take top 100.

Explainability & Data Validation


- Deterministic, fact-grounded reasoning per candidate — cites real years, title, named skills, employers, signal values, and names honest gaps (e.g., '120-day notice; Toronto-based — relocation risk').
- Zero hallucination: every term traces to the candidate's own profile (enforced by tests). Variation by dominant factor; tone tracks the rank band.
- Honeypot/anomaly detector (explainable rules): ≥ 3 advanced/expert skills with 0 months, date contradictions, tenure overflowing experience, invalid education — flags 0.042% of the pool, no false positives on real profiles.
- Schema-tolerant parsing; sentinels (-1 github/offer, empty assessments — 60–76% of pool) treated as neutral, never penalized. Fully seeded/deterministic.

End-to-End Workflow

- Phase 1 — Offline precompute (network OK, no time limit): JD →(Claude)→ jd_rubric.json; candidates → text blobs → bge-small embeddings (fp16) + BM25; Claude → eval labels. Artifacts committed.
- Phase 2 — Rank-time (CPU, no network, < 5 min): load artifacts → score 5 components → gates × behavioral → honeypot zero → sort/tie-break → top-100 → grounded reasoning → submission.csv.
- Single reproduce command: `python rank.py --candidates ./data/candidates.jsonl --out ./submission.csv` → passes the official validator.

System Architecture

Lighthouse — offline-LLM-augmented hybrid ranker



Results & Performance

- Label-INDEPENDENT evidence (no trust in our labels needed):
 - 0 honeypots in the top-100 (audited over full 100K; DQ threshold >10%). All 100/100 hold an AI/ML/IR/DS/Search/NLP title — 0 non-technical; the sample_submission instead ranks HR/Accountants at #1–20.
 - Trap resistance: the keyword baseline puts 13/32 keyword-stuffers in its top-25; Lighthouse admits 0. Remove the anti-trap stack and honeypots climb (median rank 209→143).
 - Live demo sample (100): the actual top 80 from the 100K submission + 10 keyword-stuffers + 10 honeypots. Ranks 1-80 byte-identical to the submission (real engineers, score 0.93-1.00); ranks 81-90 = stuffers with NON-TECHNICAL badge (≈40× lower score); ranks 91-100 = honeypots, every one scoring 0.0. The live Space is an in-browser replay of the official ranking.
- Directional only (self-labeled — NOT absolute accuracy): vs keyword baseline composite 0.998 vs 0.553. Labeler & ranker share assumptions, so a ~perfect NDCG@10 = internal consistency, not validated accuracy — the gap + ablation deltas are the signal. A blind human-label harness ships for an independent check.
- Constraints met: rank step CPU-only, no network, ~45 s (well under the 5-min budget) over 100K; embeddings precomputed (~75 MB fp16); 127 tests green, lint+mypy clean.

Technologies Used

- sentence-transformers BAAI/bge-small-en-v1.5 — small, CPU-friendly, strong retrieval quality; precomputed so rank-time needs no model.
- rank_bm25 — optional offline lexical index; precomputed but never loaded at rank-time (ranking is semantic + deterministic features).
- numpy / pandas — the entire rank step: fast, deterministic, dependency-light (fits the 5-min / 16 GB / CPU box).
- Claude (offline only) — JD→rubric parsing and eval labeling; never called at rank-time, no candidate data sent to any API.
- pytest (127 tests), FastAPI + vanilla JS/anime.js animated UI (HF Spaces sandbox), python-pptx + matplotlib (this deck).

- GitHub: <https://github.com/jacklachan/lighthouse-indiaruns>
- Live sandbox (HF Spaces): <https://huggingface.co/spaces/Auenchanters/lighthouse>

Reproduce: `pip install -r requirements.txt → python precompute.py → python rank.py → python validate_submission.py submission.csv`

- Demo video: [placeholder — add link before submission]
- AI use: Claude offline for JD parsing + eval labeling + coding assistant; no candidate data hit any API at rank-time (see `submission_metadata.yaml`).



INDIA.RUNS

Build what next India runs on

THANK YOU